


Best Practices for Optimizing Docker Images



Nigel Brown

@n_brownuk www.windsock.io





Mia is investigating the adoption of Docker

- **Adept at using Dockerfiles for encapsulating apps**
- **Knows how to use containers to develop with different languages**
- **What are the best practices for authoring Dockerfiles?**

Mia wants to stand on the shoulders of giants!



Module Outline



Coming up:

- Relationship between image layers and image size
- Dockerfile instruction sequencing for an efficient workflow
- Multi-stage Dockerfiles for optimizing image size
- Putting it all together





Anatomy of an Image

Understanding how Docker images are constructed is key to managing their size.



```

{
  "Config": {
    .
    .

    "workingDir": "/app
  }

  .
  .

  "RootFS": {
    "Type": "layers",
    "Layers": [
      "sha256:5a8512b2 . . . .",
      "sha256:de6a6a91 . . . .",
      "sha256:2470436b . . . ."
    ]
  }
}

```

◀ Image configuration object


◀ Working directory

◀ Filesystem definition for derived containers

◀ Content layers that make up the filesystem

Dockerfile Instruction Types

An image build processes a sequence of Dockerfile instructions




Instructions

Dockerfile instructions define the content and nature of images



Metadata

Instructions that define how derived containers will get executed



Content

Instructions that create files and directories for the image



Content Creating Dockerfile Instructions

COPY Instruction

Used to copy content
from the build context
into the image

ADD Instruction

Like COPY instruction
but can retrieve
remote content

RUN Instruction


Executes commands
to generate additional
image content




Processing a COPY, ADD or RUN instruction adds a new content layer to the image.



Image Layers



Container Filesystem



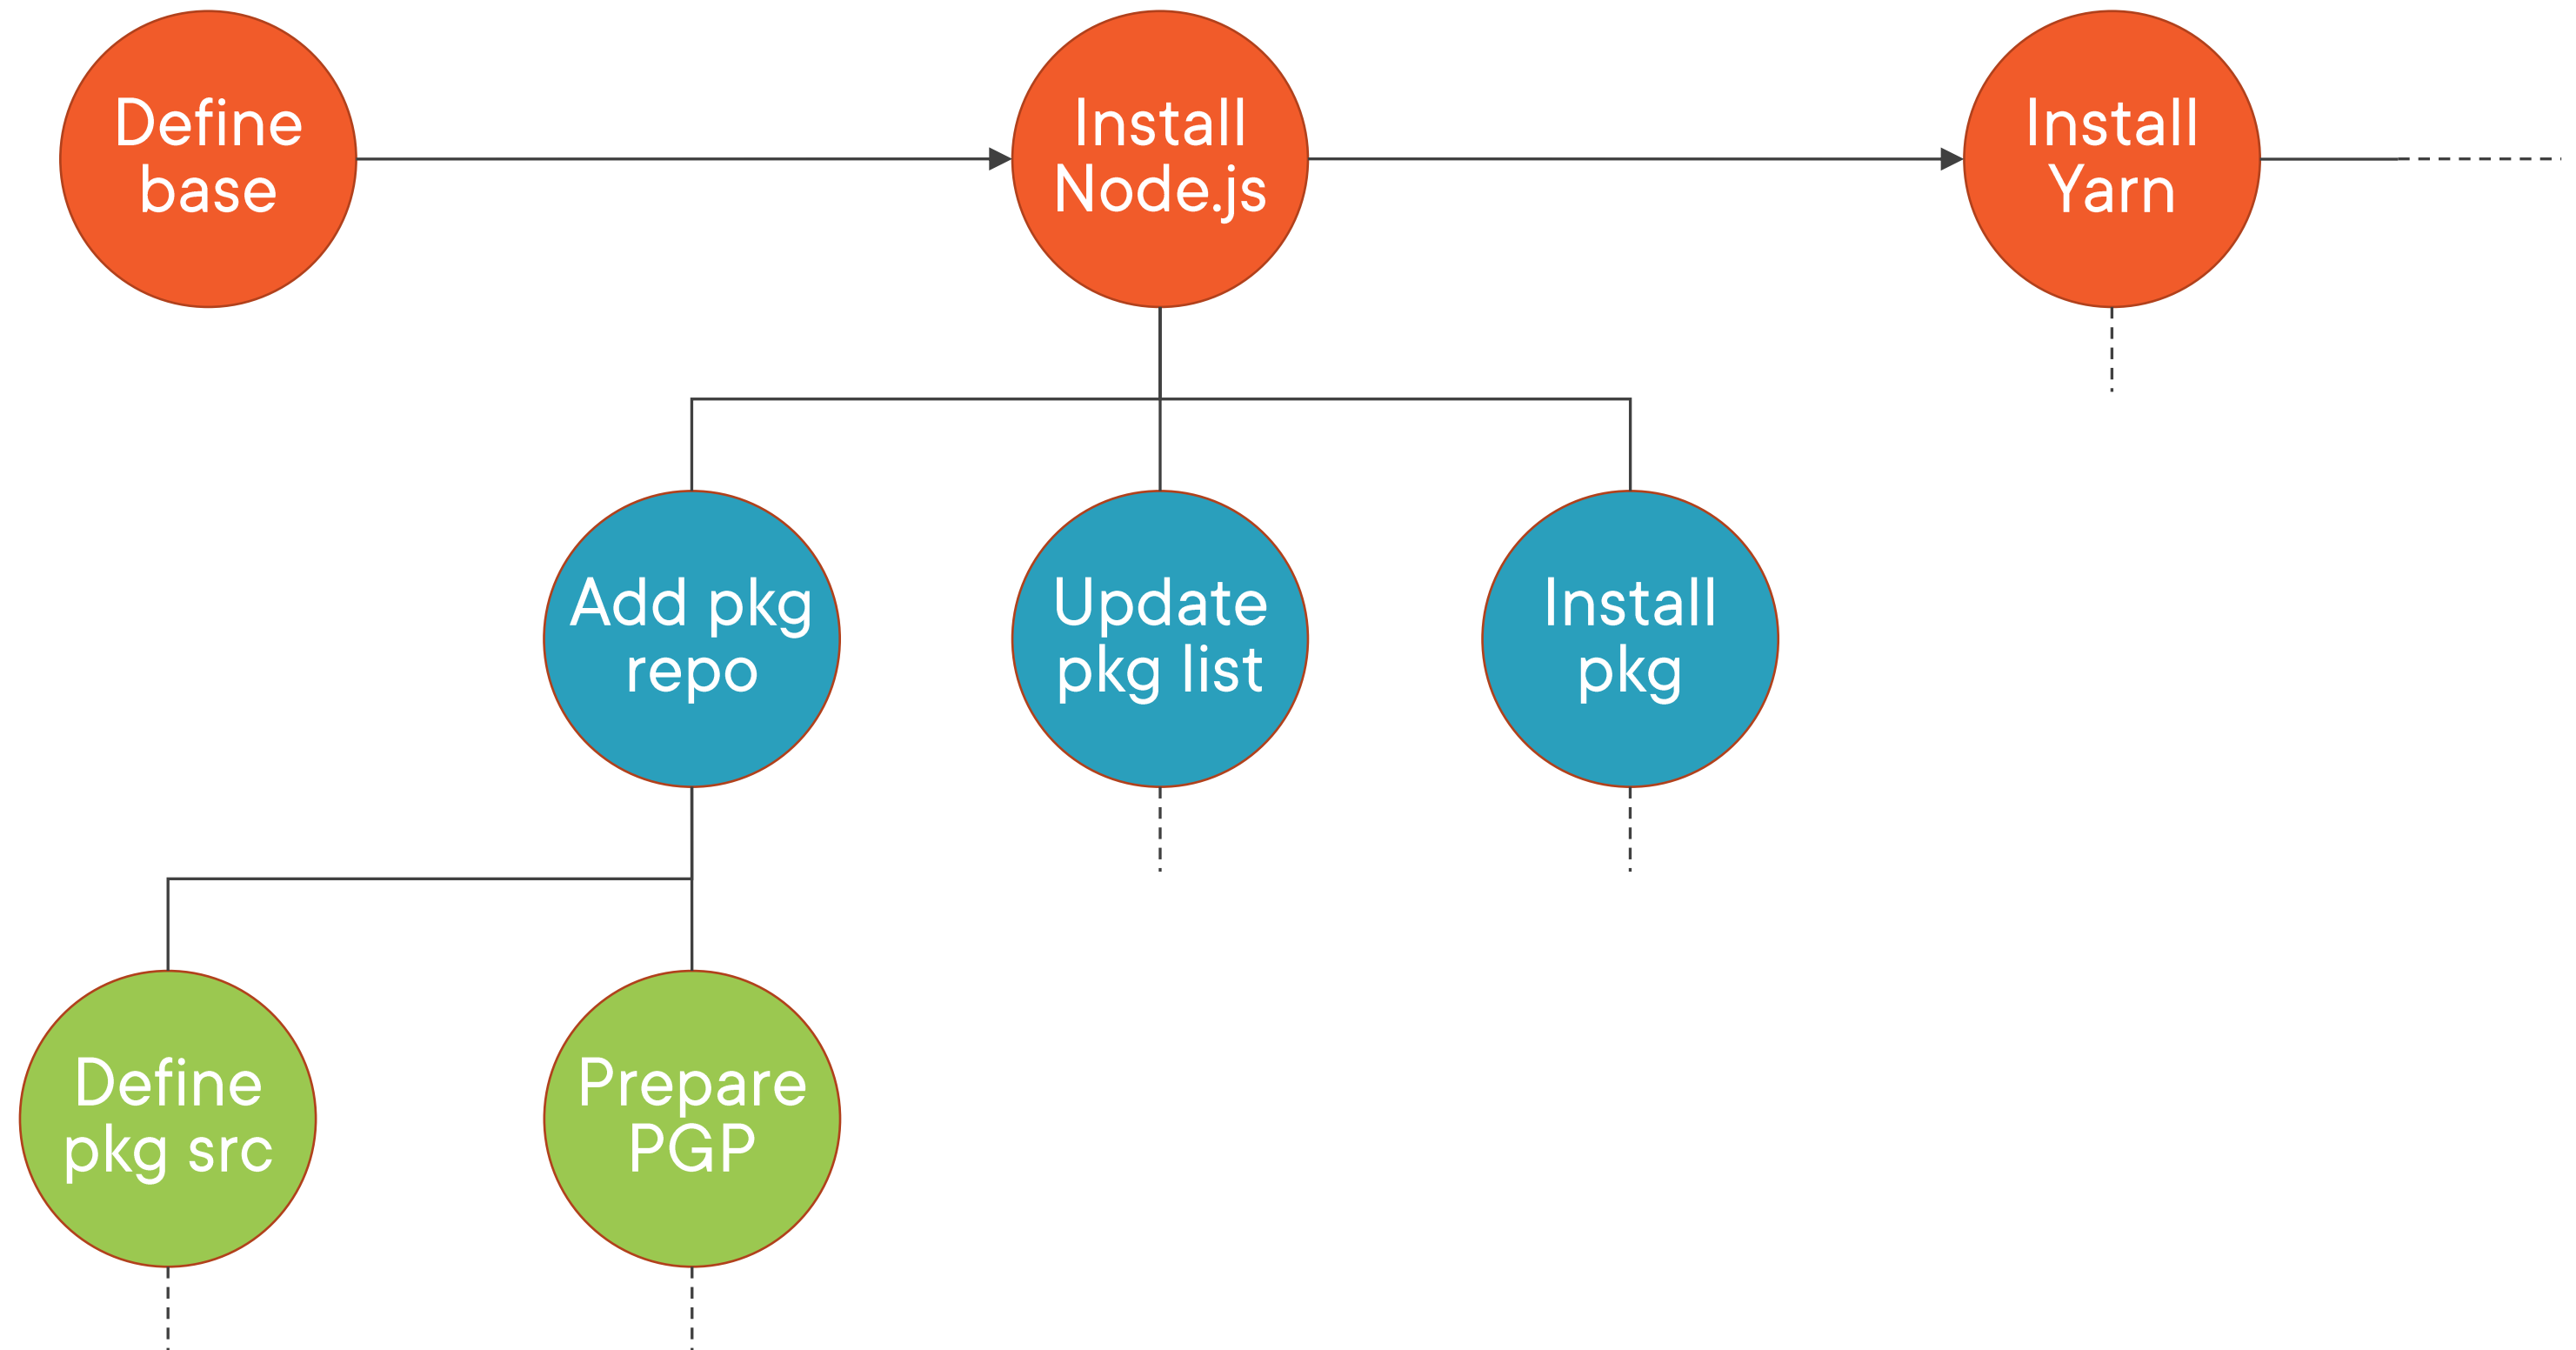


Dockerfile Scenario

The goal for our scenario is to install the Node.js runtime and the Yarn package manager into our image.



Build Steps



Side Effects

Dockerfile

```
FROM debian:buster
```

```
RUN apt-get update
```

```
RUN apt-get install -y \
    curl \
    ca-certificates \
    gnupg
```

<snip>

```
RUN apt-get update
```

```
RUN apt-get install -y \
    nodejs \
    yarn
```

<snip>



Temporary Content

The additional content is required temporarily

It makes the image larger than it needs to be

The content needs to be removed after use

Content Removal

Dockerfile

```
FROM debian:buster
```

```
RUN apt-get update
```

```
RUN apt-get install -y \
    curl \
    ca-certificates \
    gnupg
```

<snip>


```
RUN apt-get update
```

```
RUN apt-get install -y \
    nodejs \
    yarn
```

```
RUN apt-get purge -y curl ca-certificates gnupg
```

<snip>

Hidden Content in Layers



RUN instruction to remove content

RUN instruction to add content



Logical AND Operator Use

Dockerfile

```
FROM debian:buster
```

```
RUN apt-get update    && \
```

```
RUN apt-get install -y \
    curl \
    ca-certificates \
    gnupg    && \
```

```
<snip>    && \
```

```
RUN apt-get update    && \
```

```
RUN apt-get install -y \
    nodejs \
    yarn    && \
```

```
RUN apt-get purge -y curl gnupg
```

```
<snip>
```



Build Cache


Docker uses a local cache of image build steps.
Careful placement of Dockerfile instructions can
maximize cache hits.




Image Chain



Each Dockerfile instruction processed during a build results in the creation of an intermediary image that is part of the build cache



These images are created by ‘committing’ containers created from the image associated with the preceding Dockerfile instruction



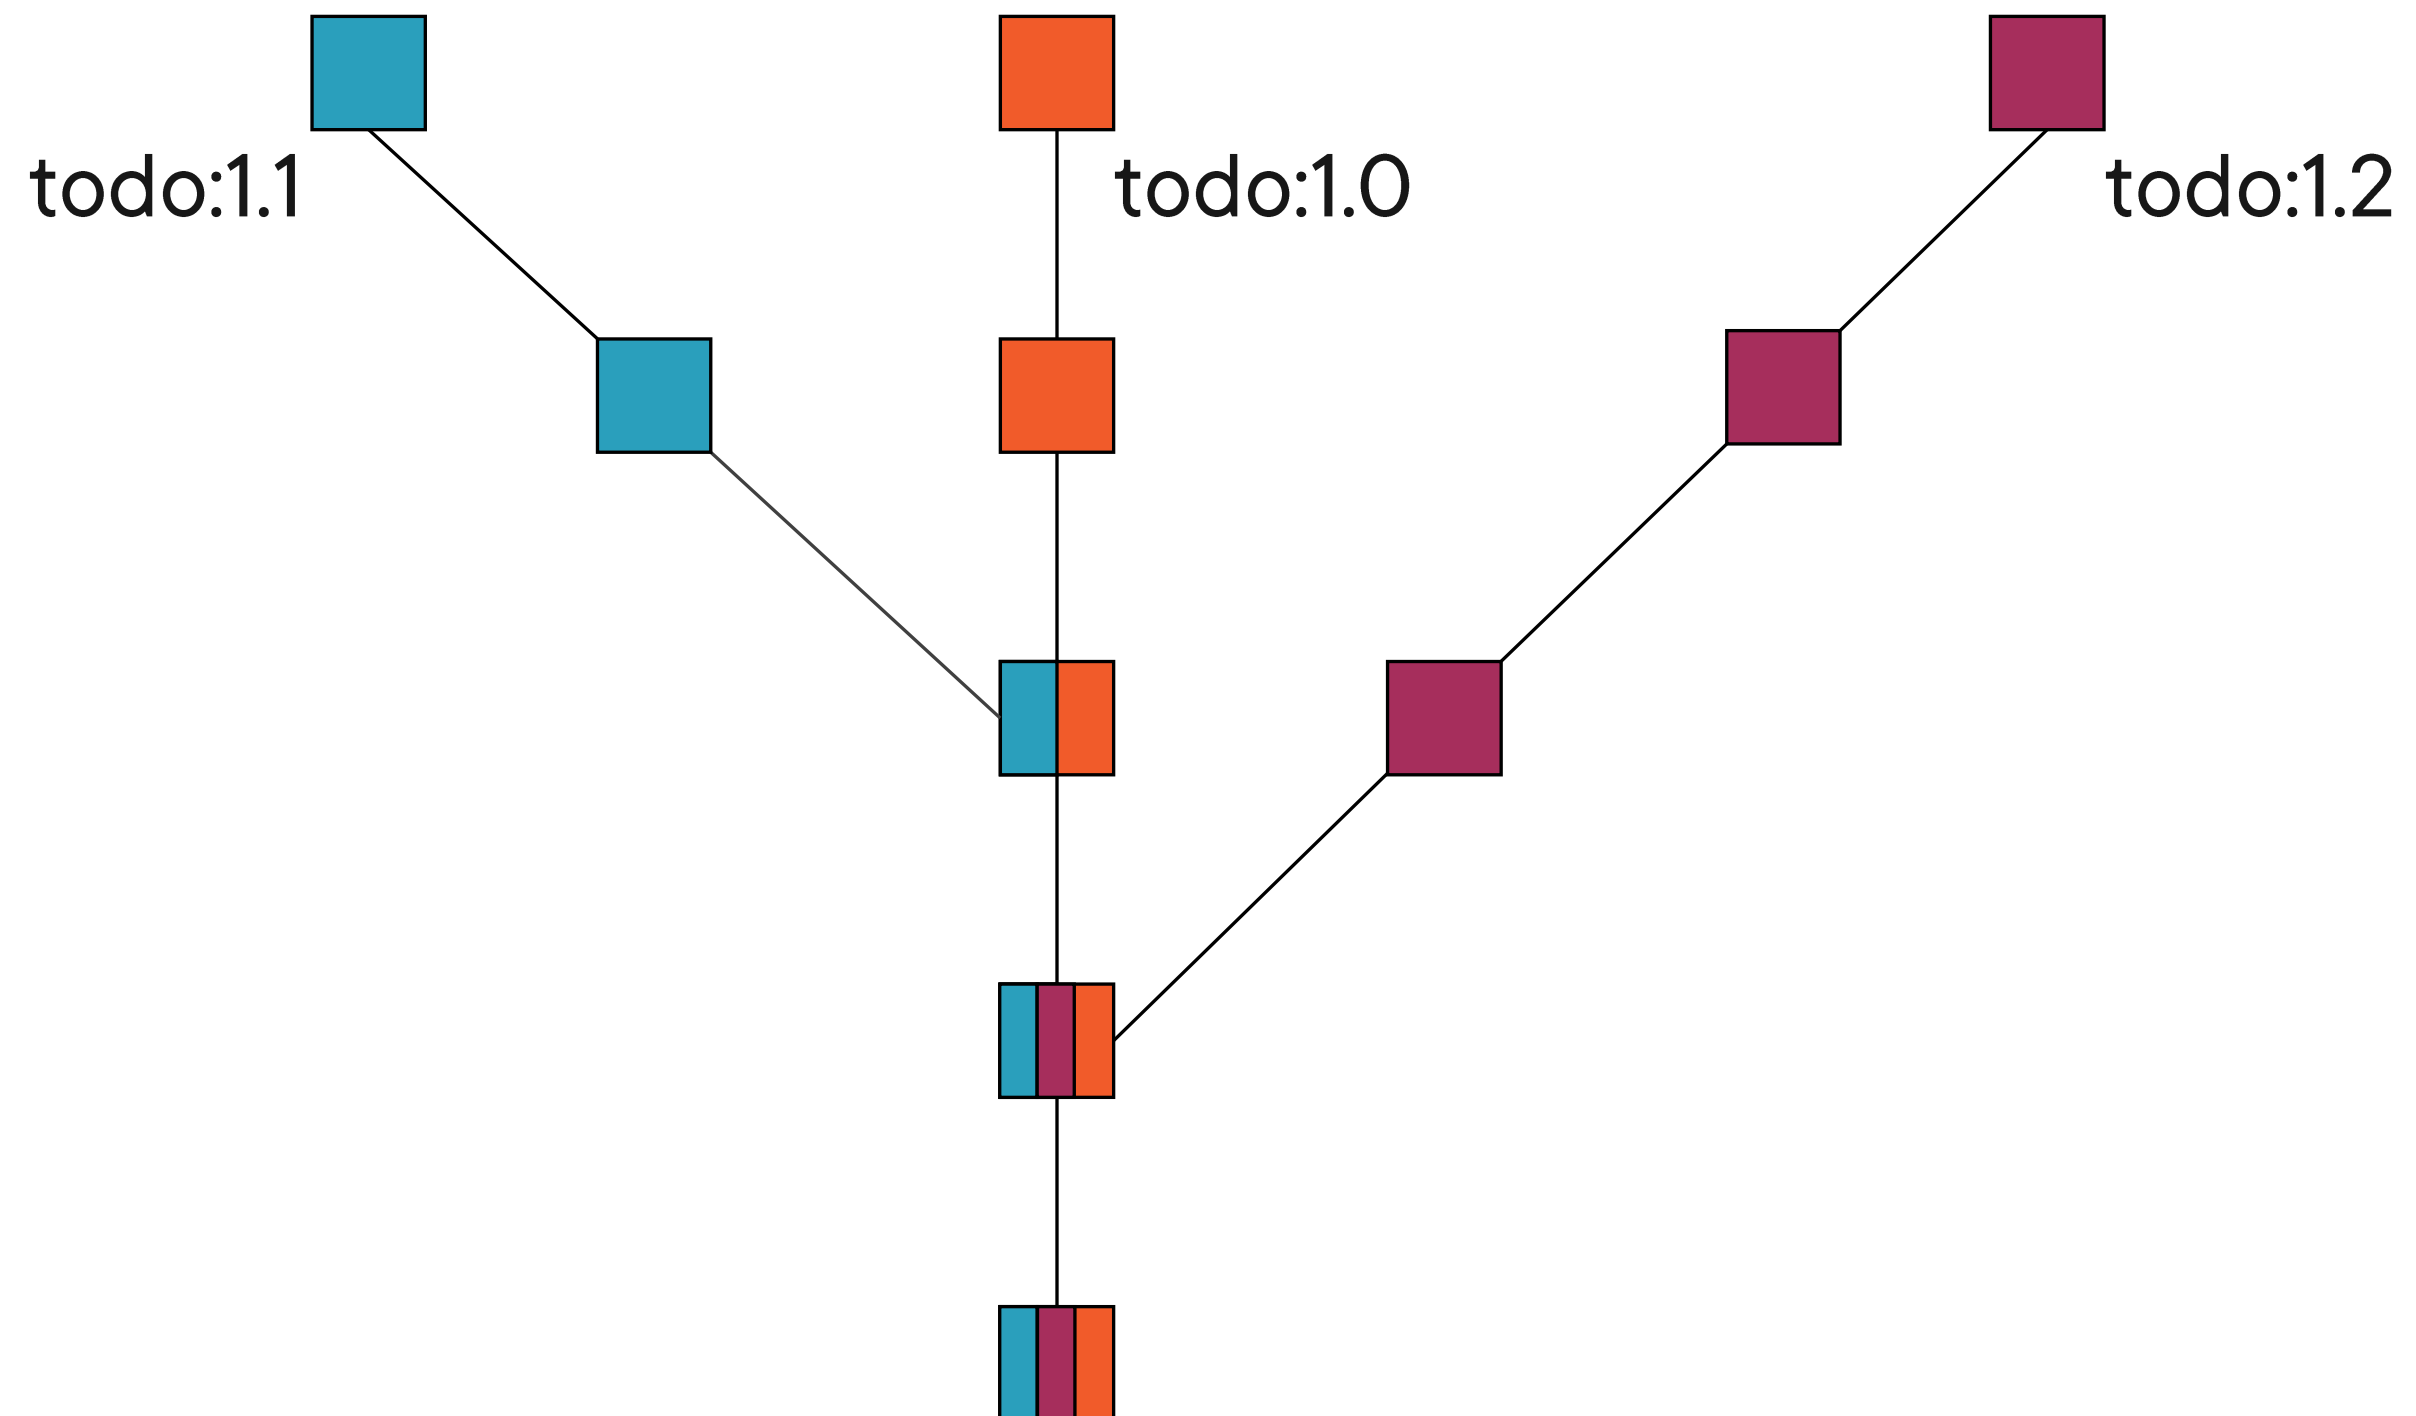
Images reference their parent image and thereby create an implicit chain of images that represent a sequence of instructions



Docker will pass over instructions that form a sequence that already exists in the build cache.



Using the Build Cache



Hitting and Missing



Instruction change
Adding, removing or
altering an instruction
invalidates the cache



Checksum check
Content change in
build context will
invalidate the cache



Command output
Consequences of
command execution
are not checked

Sequencing Dockerfile Instructions

Before

<snip>

WORKDIR /app

COPY . .

RUN yarn install

<snip>

After

<snip>

WORKDIR /app


COPY package.json yarn.lock ./

RUN yarn install

COPY spec src ./

<snip>


The Practicalities



Analyze the dependencies between Dockerfile instructions to determine ordering constraints



Order Dockerfile instructions according to the frequency of change; less frequent first, more frequent last



Where it's beneficial, split COPY Dockerfile instructions that copy content from the build context



Multi-stage Dockerfiles can play a bit part in optimizing the size of images.



Dockerfile

```
FROM debian:buster
```

```
RUN apt-get update                                && \
    apt-get install -y --no-install-recommends curl ca-certificates gnupg && \
    curl -s https://deb.nodesource.com/gpgkey/nodesource.gpg.key |      \
        apt-key add -                                                    && \
    echo 'deb https://deb.nodesource.com/node_14.x buster main' |        \
        tee /etc/apt/sources.list.d/nodesource.list                     && \
    curl -sS https://dl.yarnpkg.com/debian/pubkey.gpg | apt-key add -    && \
    echo "deb https://dl.yarnpkg.com/debian/ stable main" |              \
        tee /etc/apt/sources.list.d/yarn.list                            && \
    apt-get update                                                        && \
    apt-get install -y --no-install-recommends nodejs yarn              && \
    apt-get purge -y curl gnupg                                         && \
    rm -rf /var/lib/apt/lists/*
```

Dockerfile

```
FROM debian:buster
```

```
RUN apt-get update                                && \  
    apt-get install -y --no-install-recommends curl ca-certificates gnupg && \  
    curl -s https://deb.nodesource.com/gpgkey/nodesource.gpg.key |          \  
        apt-key add -                                && \  
    echo 'deb https://deb.nodesource.com/node_15.x buster main' |            \  
        tee /etc/apt/sources.list.d/nodesource.list                        && \  
    curl -sS https://dl.yarnpkg.com/debian/pubkey.gpg | apt-key add -      && \  
    echo "deb https://dl.yarnpkg.com/debian/ stable main" |                 \  
        tee /etc/apt/sources.list.d/yarn.list                              && \  
    apt-get update                                && \  
    apt-get install -y --no-install-recommends nodejs yarn                  && \  
    apt-get purge -y curl gnupg                                             && \  
    rm -rf /var/lib/apt/lists/*
```

Dockerfile

```
FROM debian:buster AS base
```

```
RUN apt-get update
```

```
RUN apt-get install -y --no-install-recommends curl ca-certificates gnupg
```

```
RUN curl -s https://deb.nodesource.com/gpgkey/nodesource.gpg.key | \
    apt-key add -
```

```
RUN echo 'deb https://deb.nodesource.com/node_14.x buster main' | \
    tee /etc/apt/sources.list.d/nodesource.list
```

```
RUN curl -sS https://dl.yarnpkg.com/debian/pubkey.gpg | apt-key add -
```

```
RUN echo "deb https://dl.yarnpkg.com/debian/ stable main" | \
    tee /etc/apt/sources.list.d/yarn.list
```


```
RUN apt-get update
```

```
RUN apt-get install -y --no-install-recommends nodejs yarn
```

```
RUN apt-get purge -y curl gnupg
```

```
RUN rm -rf /var/lib/apt/lists/*
```

Profiting from Multi-stage Dockerfiles



Return to a RUN instruction for each command



Maximizes the use of the build cache



Temporary content resides in a previous stage

Choice is a trade-off

- **Image size vs build speed**
- **Be wary when sharing the build cache**



Demo



Creating an optimal image build for an application

- Start with a sub-optimal Dockerfile
- Minimize content using a single layer
- See the effect of careful Dockerfile instruction sequencing
- Optimize image size using multiple stages



Up Next:

Making Configuration Data Available to
Containerized Applications



Module Summary



What we covered:

- Relationship between layers and size
- Concatenating commands using the AND operator
- Docker instruction sequencing
- Enhanced builds with multi-stage Dockerfiles

